

# **SCOM/ SRSI**

# **Application Layer, Web**

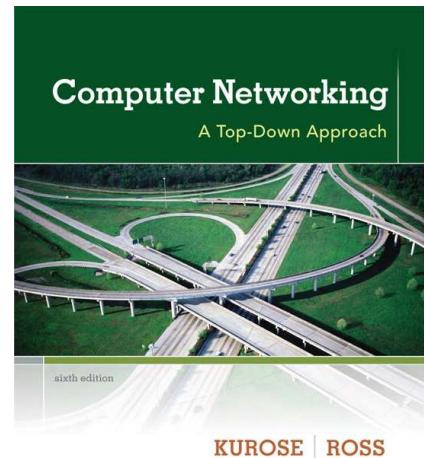
Ana Aguiar

DEEC, FEUP

2018-19

# Chapter 2

## Application Layer



### A note on the use of these ppt slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- ❖ If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- ❖ If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2012  
© J.F Kurose and K.W. Ross, All Rights Reserved

**Computer  
Networking: A Top  
Down Approach**  
**6<sup>th</sup> edition**  
**Jim Kurose, Keith Ross**  
**Addison-Wesley**  
**March 2012**

# Chapter 2: outline

**2.1 principles of network applications**

**2.2 Web and HTTP**

**2.3 FTP**

**2.4 electronic mail**

– SMTP, POP3, IMAP

**2.5 DNS**

**2.6 P2P applications**

**2.7 socket programming with UDP and TCP**

# Chapter 2: application layer

## our goals:

- conceptual, implementation aspects of network application protocols
  - transport-layer service models
  - client-server paradigm
  - peer-to-peer paradigm
- learn about protocols by examining popular application-level protocols
  - HTTP
  - FTP
  - SMTP / POP3 / IMAP
  - DNS
- creating network applications
  - socket API

# Some network apps

- e-mail
- web
- text messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video  
(YouTube, Hulu, Netflix)
- voice over IP (e.g., Skype)
- real-time video conferencing
- social networking
- search
- ...
- ...

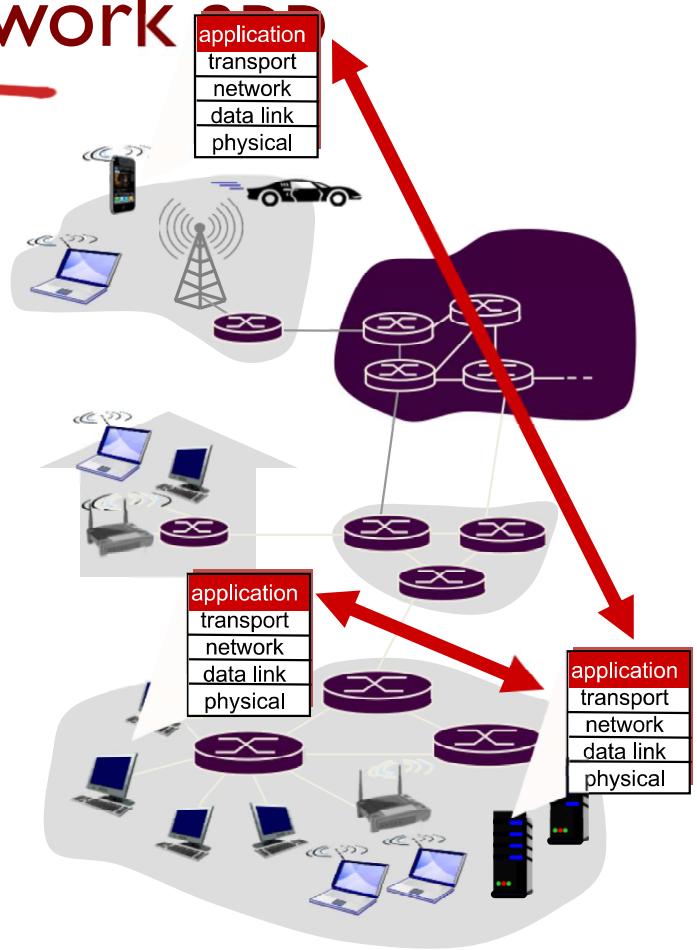
# Creating a network

write programs that:

- run on (different) end systems
- communicate over network
- e.g., web server software communicates with browser software

no need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation

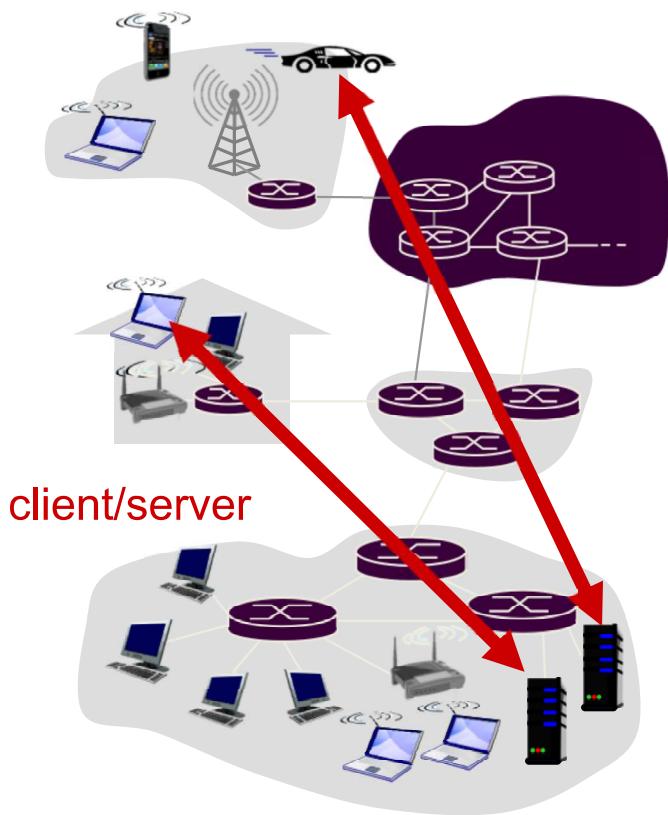


# Application architectures

possible structure of applications:

- client-server
- peer-to-peer (P2P)

# Client-server architecture



## server:

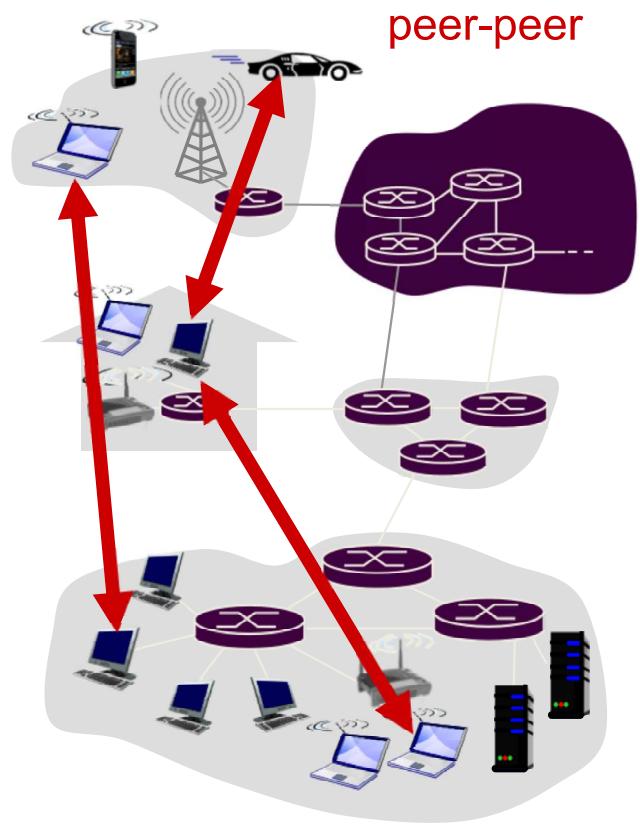
- always-on host
- permanent IP address
- data centers for scaling

## clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other

# P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
  - **self scalability** – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
  - complex management



# Processes communicating

**process:** program running within a host

- within same host, two processes communicate using **inter-process communication** (defined by OS)
- processes in different hosts communicate by exchanging **messages**

clients, servers

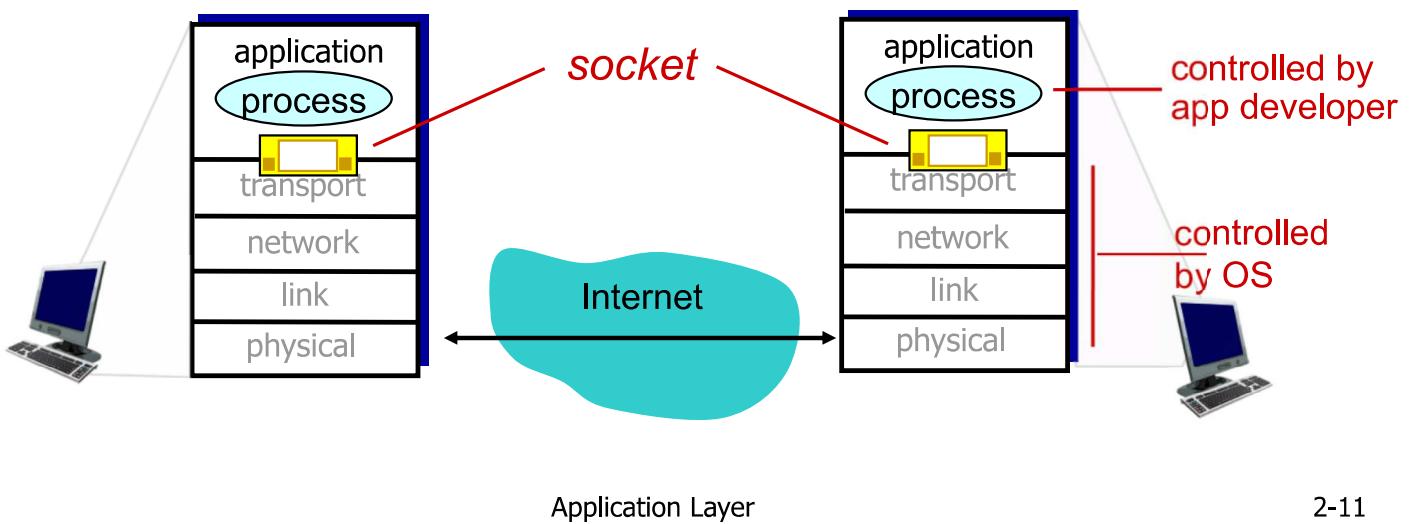
**client process:** process that initiates communication

**server process:** process that waits to be contacted

- ❖ aside: applications with P2P architectures have client processes & server processes

# Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
  - sending process shoves message out door
  - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



# Addressing processes

- to receive messages, process must have *identifier*
- host device has unique 32-bit IP address
- **Q:** does IP address of host on which process runs suffice for identifying the process?
  - **A:** no, *many* processes can be running on same host
- *identifier* includes both IP address and port numbers associated with process on host.
- example port numbers:
  - HTTP server: 80
  - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
  - IP address: 128.119.245.12
  - port number: 80
- more shortly...

# App-layer protocol defines

- types of messages exchanged,
  - e.g., request, response
- message syntax:
  - what fields in messages & how fields are delineated
- message semantics
  - meaning of information in fields
- rules for when and how processes send & respond to messages

open protocols:

- defined in RFCs
- allows for interoperability
- e.g., HTTP, SMTP

proprietary protocols:

- e.g., Skype

# What transport service does an app need?

## **data integrity**

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

## **timing**

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

## **throughput**

- ❖ some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- ❖ other apps (“elastic apps”) make use of whatever throughput they get

## **security**

- ❖ encryption, data integrity, ...

## Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video:10kbps-5Mbps	yes, 100's msec
stored audio/video	loss-tolerant	same as above	
interactive games	loss-tolerant	few kbps up	yes, few secs
text messaging	no loss	elastic	yes, 100's msec yes and no

# Internet transport protocols services

## TCP service:

- *reliable transport* between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum throughput guarantee, security
- *connection-oriented*: setup required between client and server processes

## UDP service:

- *unreliable data transfer* between sending and receiving process
- *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Q: why bother? Why is there a UDP?

# Internet apps: application, transport protocols

application	application layer protocol	underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	TCP or UDP

# **WEB**

Is the web the same as the Internet?

What makes a machine part of the web?

# World Wide Web

- How does web browsing work?
- Which is the communication model?
- How do you find contents?
- Which protocols are involved?
- How could you improve the quality of experience?

# Exercise: HTTP Protocol

- Start a wireshark capture
- Open the browser
- Access [www.cloud.futurecities.up.pt](http://www.cloud.futurecities.up.pt), [www.w3c.org](http://www.w3c.org), ...
- Stop the wireshark capture
- Analyse the log
  - Find the DNS query packets
  - How much time did it take to resolve the name?
  - How many TCP connections were opened?
  - Which information is inside the query and response headers?
  - What is the meaning of that information?

# HTTP

# WWW Architectural Components

- Client, e.g. browser
- Server
- HTTP protocol
- HTML
- URL
- Proxies

# HyperText Transfer Protocol

- HyperText Transfer Protocol (HTTP): RFC 2616
  - Client-server protocol
  - Request-response
  - Evoked by the browser on the client machine and by the server process on the server machine
  - Uses TCP
  - Stateless
    - Server and client do not keep state (information) about connection

# HTTP Message

```
START LINE <CRLF>
MESSAGE_HEADER <CRLF>
<CRLF>
MESSAGE BODY <CRLF>
```

- Start Line says whether it is a request or response
- Message Header specifies options and parameters
  - <CRLF> acts as delimiter
- Message Body contains the message

# HTTP Requests

- GET: retrieve resource contents
- POST: write to resource
- PUT: update resource with contents
- DELETE: remove resource contents
- OPTIONS: retrieve information about communication capabilities and/ or content options
- HEAD: retrieve header only
- TRACE: application-level loopback (for debug)
- CONNECT: for use with proxy

# HTTP Responses

- Start line contains
  - HTTP version
  - Response code
    - 1xx            Informational
    - 2xx            Success
    - 3xx            Redirection
    - 4xx            Client Error
    - 5xx            Server Error
  - String explaining response (optional)
- How are non-text contents exchanged?
  - MIME

# Cookies

- Why are cookies needed?
  - Recall that HTTP is a stateless protocol
  - How to keep session information, e.g. shopping cart?
- What are cookies?
  - Pieces of data sent as an HTTP response
  - Stored by client (browser) and sent in subsequent interactions
  - Used by the server to keep session state
- IETF Standard RFC 6265
  - <http://tools.ietf.org/html/rfc6265>
- Are cookies dangerous?
  - <https://www.addedbytes.com/blog/are-cookies-dangerous/>



# User-server state: cookies

many Web sites use cookies

*four components:*

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

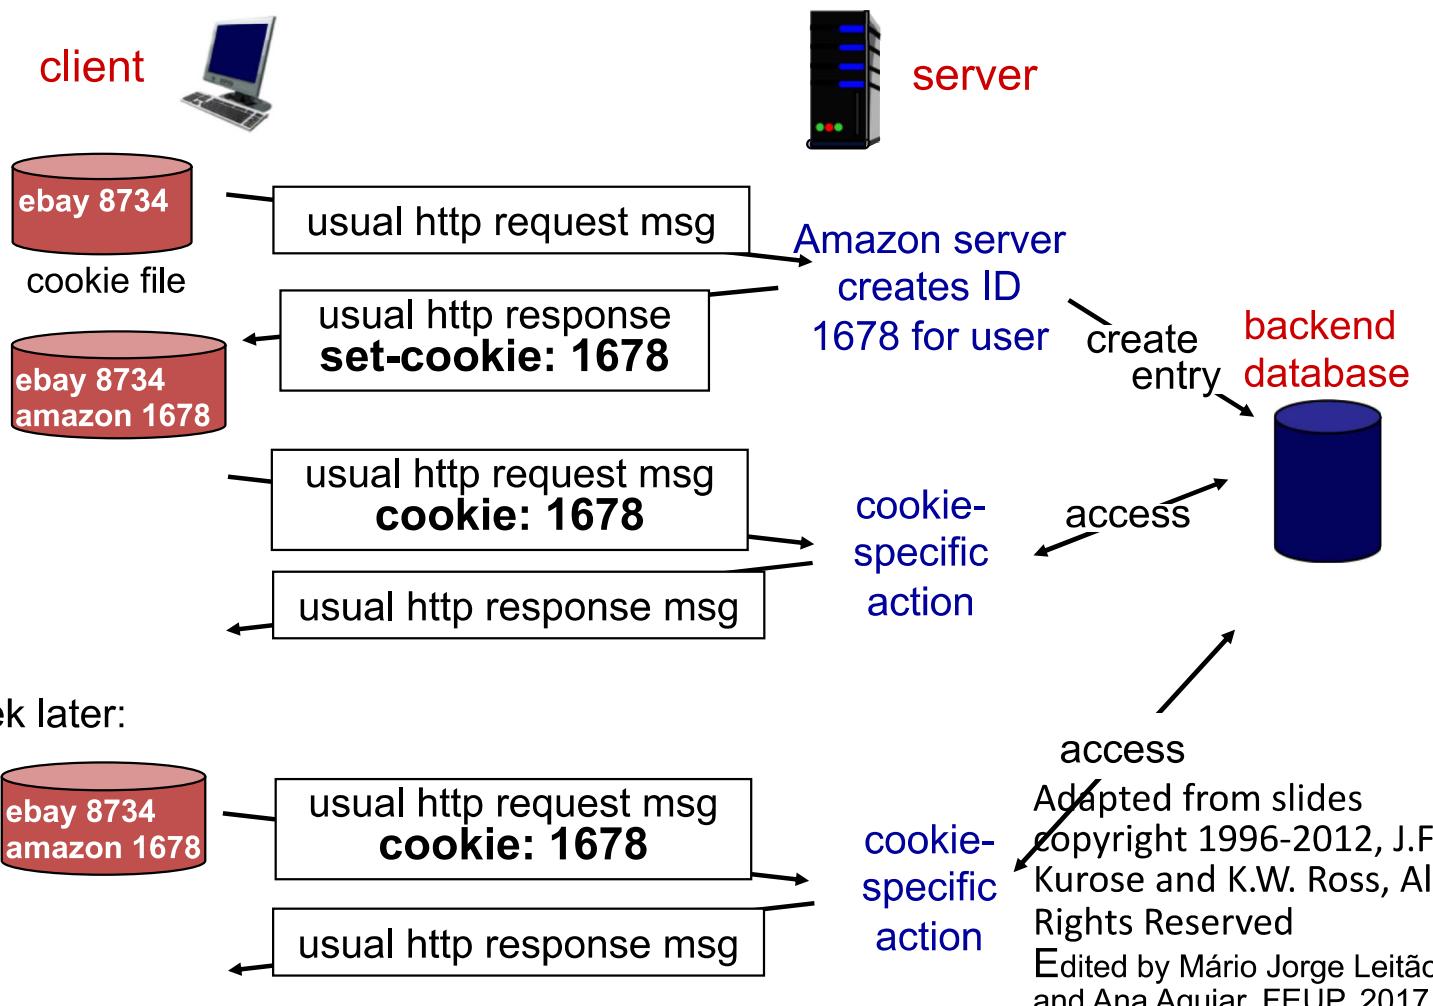
*example:*

- Susan always access Internet from PC
- visits specific e-commerce site for first time
- when initial HTTP request arrives at site, site creates:
  - unique ID
  - entry in backend database for ID

Adapted from slides copyright 1996-2012, J.F Kurose and K.W. Ross, All Rights Reserved

Edited by Mário Jorge Leitão and Ana Aguiar, FEUP, 2017

# Cookies: keeping “state” (cont.)



# Cookies (continued)

*what cookies can be used for:*

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

*aside*  
*cookies and privacy:*

- ❖ cookies permit sites to learn a lot about you
- ❖ you may supply name and e-mail to sites

*how to keep “state”:*

- ❖ protocol endpoints: maintain state at sender/receiver over multiple transactions
- ❖ cookies: http messages carry state

Adapted from slides copyright 1996-2012, J.F Kurose and K.W. Ross, All Rights Reserved  
Edited by Mário Jorge Leitão and Ana Aguiar, FEUP, 2017

# Cookies

- How can you be tracked using cookies?

# HTTP Performance

- Which metric would you use to evaluate HTTP?
- Go to the wireshark log you captured earlier today:
  - What was the page loading delay, i.e. the delay perceived by the user?
    - Assume that the browser did not add any delay
  - How many elements did the webpage have?
- What does the performance depend on?

# HTTP Performance

- What does it depend on?
  - Network performance: bandwidth, latency
  - Server performance
  - Contents
    - Amount and size of embedded elements (images, video, etc)
    - External fonts, scripts, etc
    - Further reading on website performance: Patrick Meenan. How fast is your website?. Communications of the ACM 56, 4 (April 2013), 49-55. DOI=10.1145/2436256.2436270
  - For the user: also on processing and rendering time

# HTTP connections

## *non-persistent HTTP*

- at most one object sent over TCP connection
  - connection then closed
- downloading multiple objects required multiple connections

## *persistent HTTP*

- multiple objects can be sent over single TCP connection between client, server

Adapted from slides copyright 1996-2012, J.F Kurose and K.W. Ross, All Rights Reserved  
Edited by Mário Jorge Leitão and Ana Aguiar, FEUP, 2017

# Non-persistent HTTP

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

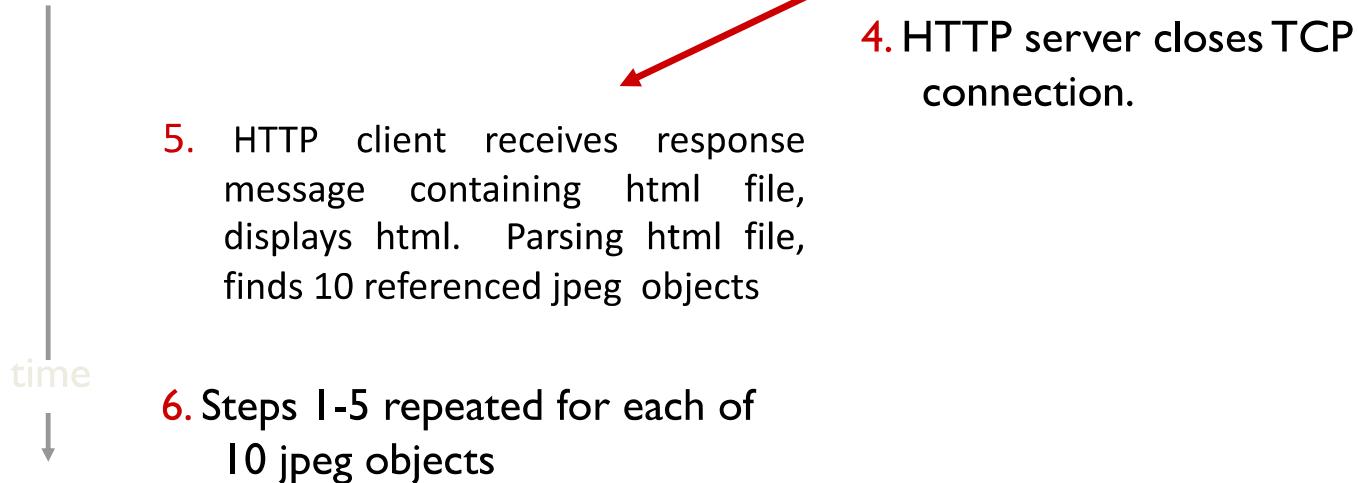
(contains text,  
references to 10  
jpeg images)

- 1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80
- 1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client
2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`
3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time

Adapted from slides copyright 1996-2012, J.F Kurose and K.W. Ross, All Rights Reserved  
Edited by Mário Jorge Leitão and Ana Aguiar, FEUP, 2017

# Non-persistent HTTP (cont.)



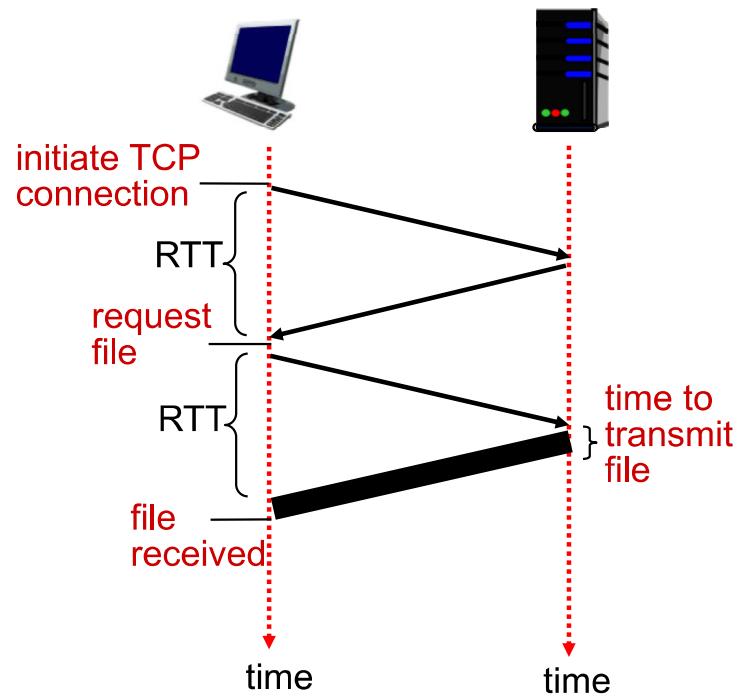
Adapted from slides copyright 1996-2012, J.F Kurose and K.W. Ross, All Rights Reserved  
Edited by Mário Jorge Leitão and Ana Aguiar, FEUP, 2017

# Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =  
$$2\text{RTT} + \text{file transmission time}$$



Adapted from slides copyright 1996-2012, J.F Kurose and K.W. Ross, All Rights Reserved  
Edited by Mário Jorge Leitão and Ana Aguiar, FEUP, 2017

# Persistent HTTP

*non-persistent*

*issues:*

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

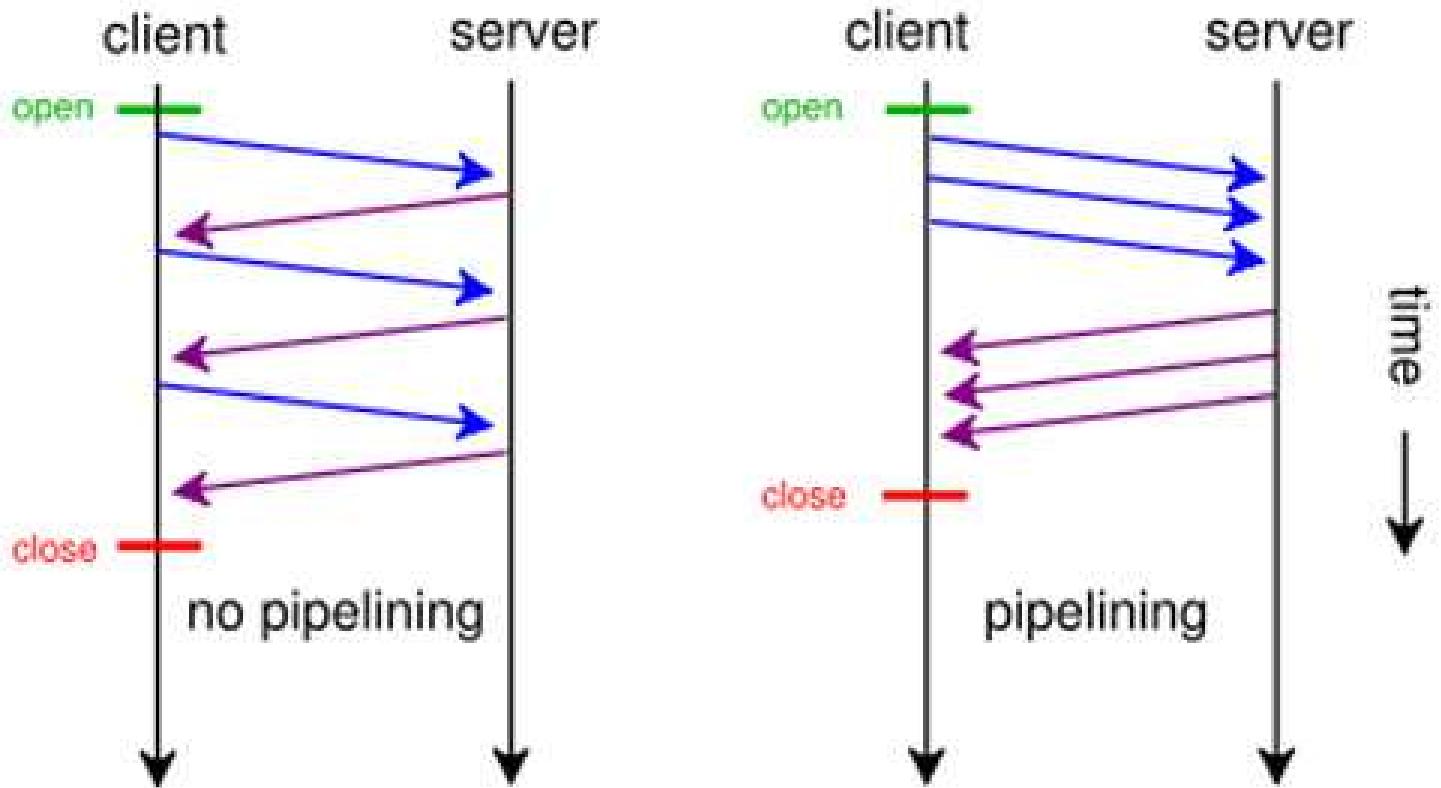
*HTTP*

*persistent HTTP:*

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

Adapted from slides copyright 1996-2012, J.F Kurose and K.W. Ross, All Rights Reserved  
Edited by Mário Jorge Leitão and Ana Aguiar, FEUP, 2017

# HTTP/1.1 Pipelining



Source: <http://slides.com/ipeychev/http-2-0-and-quic-protocols-of-the-near-future-and-why-they-re-important#/>

# Homework

Admita pretende aceder a uma página web que contém um código html base e referências a 5 objetos (imagens) que são descarregados para apresentar a totalidade da página. Assuma que RTT = 40 ms e despreze os atrasos de transmissão face a RTT.

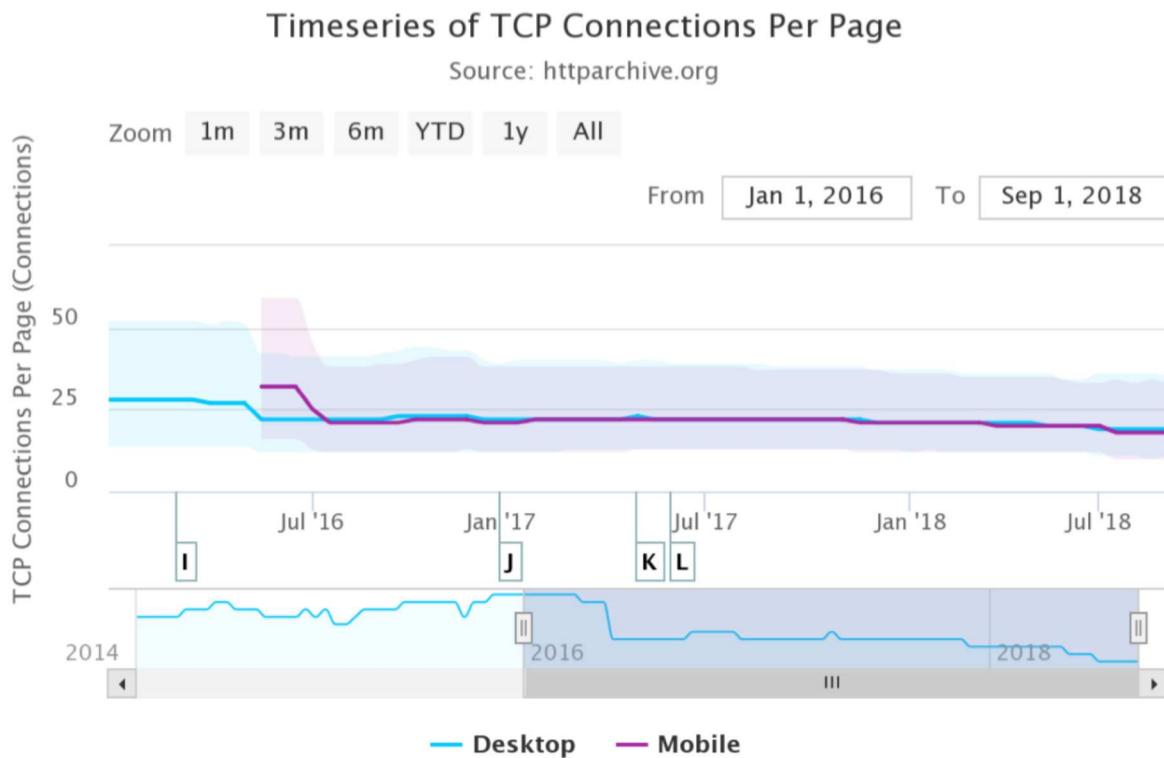
Quanto tempo é necessário para receber a página no caso de:

- a) HTTP não persistente sem conexões TCP paralelas.
- b) HTTP não persistente com 5 conexões TCP paralelas.
- c) HTTP persistente com pedidos imediatos dos 5 objetos logo que sejam identificados na página, sem aguardar a respetiva receção (em pipelining - com sobreposição de transações).

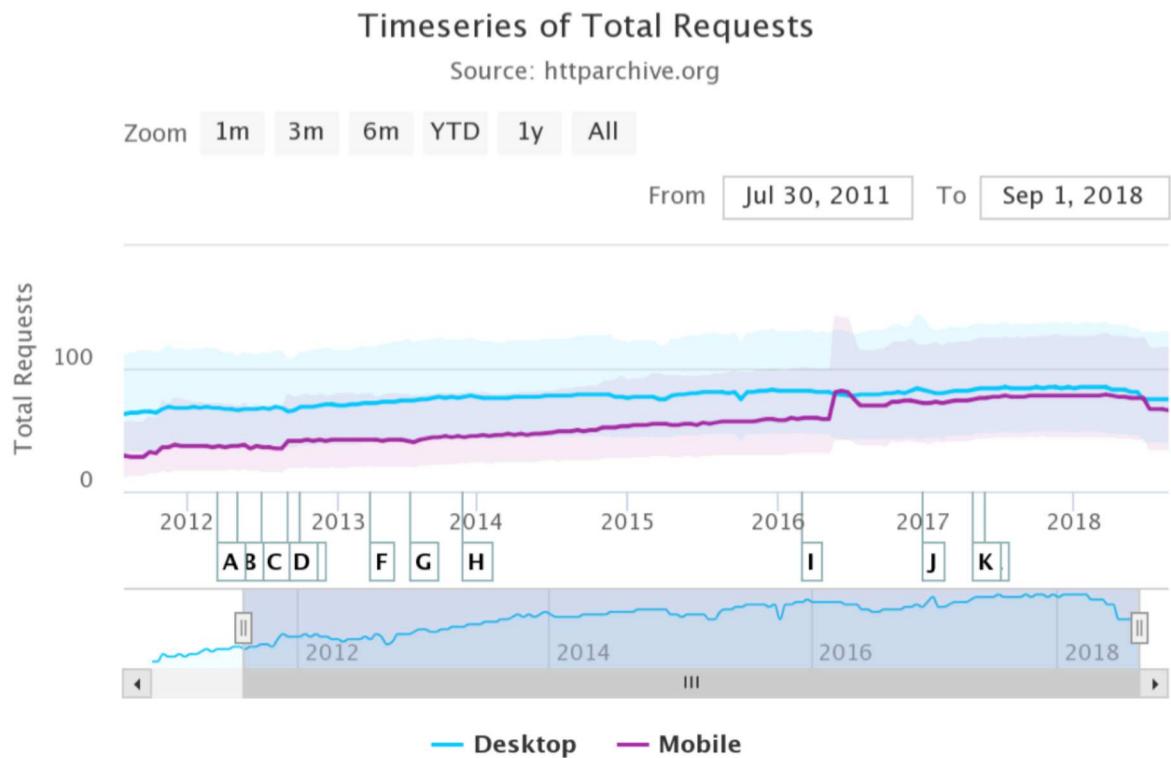
# TCP Connections Summary

- HTTP 1.0
  - Non-persistent TCP connections
- HTTP 1.1
  - Persistent TCP connection
    - Lower connection setup overhead
    - Take advantage of expanded congestion window
  - Request pipelining
  - Improved caching support with new headers
    - We will see this next week
  - Others...

# TCP Connections per Page



# Objects per Page



# Web Application Today

- > 90 resources
- From > 15 distinct hosts
- > 1.3MB compressed data
- HTTP requests are large because of header
- Most TCP flows are small
- What does this mean for response time?
- How to reduce latency?
  - HTTP/2
  - Content Delivery Networks (next week)
  - Both
  - Improve transport

# **HTTP/2.0**

# HTTP/2.0: Main Features

- Binary
  - To reduce overhead, on the wire and processing
- Full multiplexing of requests and responses
  - To reduce latency
- Parallelism on single TCP connection
  - To reduce network and system overhead
- Server push
  - To probabilistically reduce latency
- Header compression
  - To reduce overhead

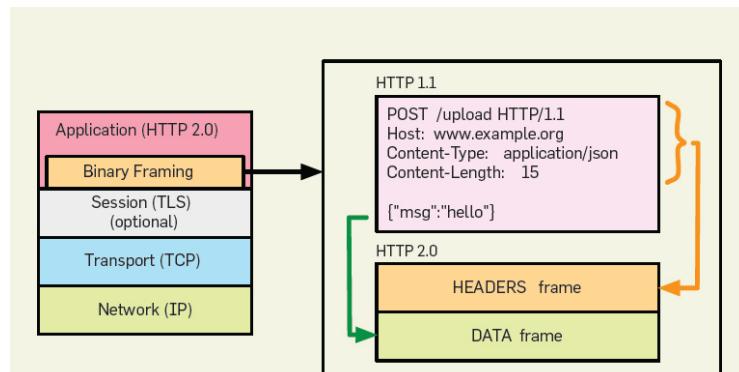
<http://cacm.acm.org/magazines/2013/12/169944-making-the-web-faster-with-http-2-0/fulltext>  
<https://http2.github.io/>

# HTTP/2.0 Full Req-Rep Multiplexing

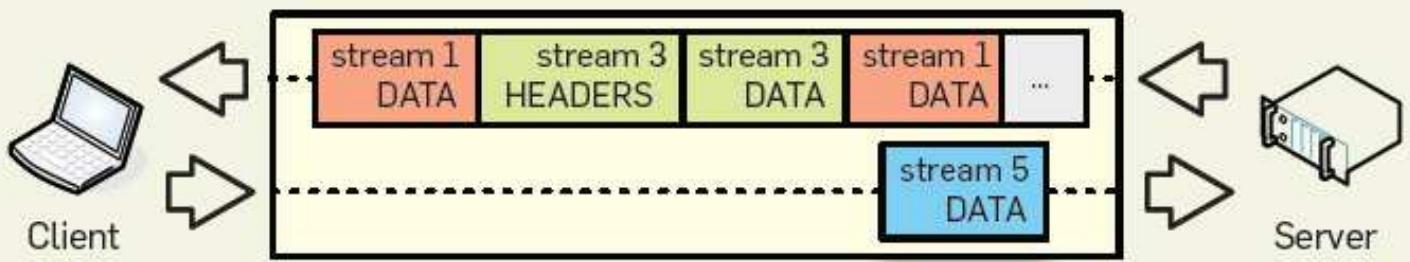
- Full request and response multiplexing
  - Stream: bi-directional byte flow
  - Message: conceptual messages that aggregates frames, e.g. a request or a response
  - Frames: basic communication unit
- Enables
  - Interleave multiple requests in parallel without blocking on any one.
  - Interleave multiple responses in parallel without blocking on any one.
  - Use a single connection to deliver many requests and responses in parallel.
  - Remove unnecessary HTTP 1.x workarounds from application code.

Source (also for images): <http://cacm.acm.org/magazines/2013/12/169944-making-the-web-faster-with-http-2-0/fulltext>

# HTTP/2.0

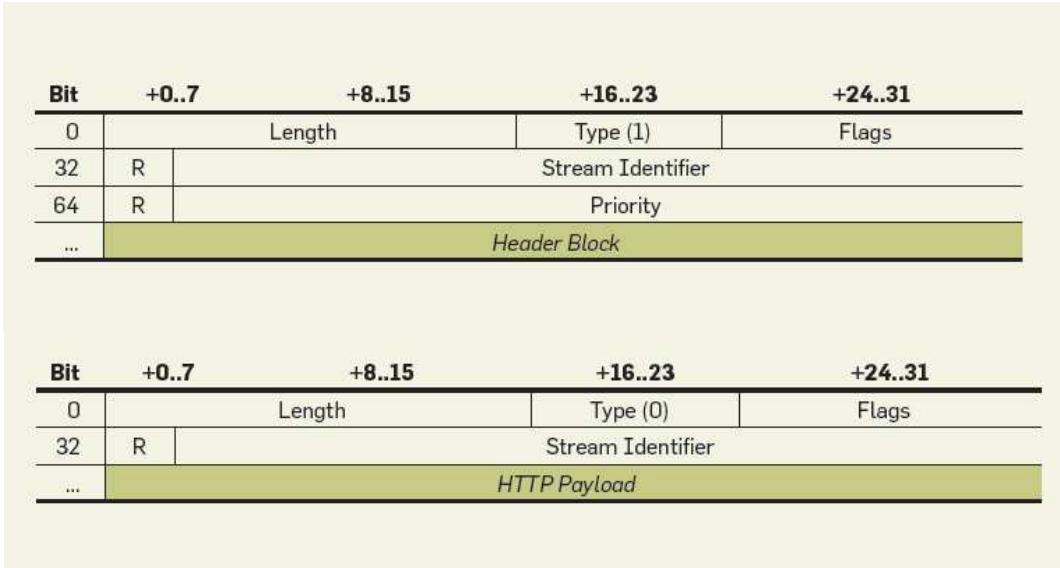


HTTP 2.0 connection



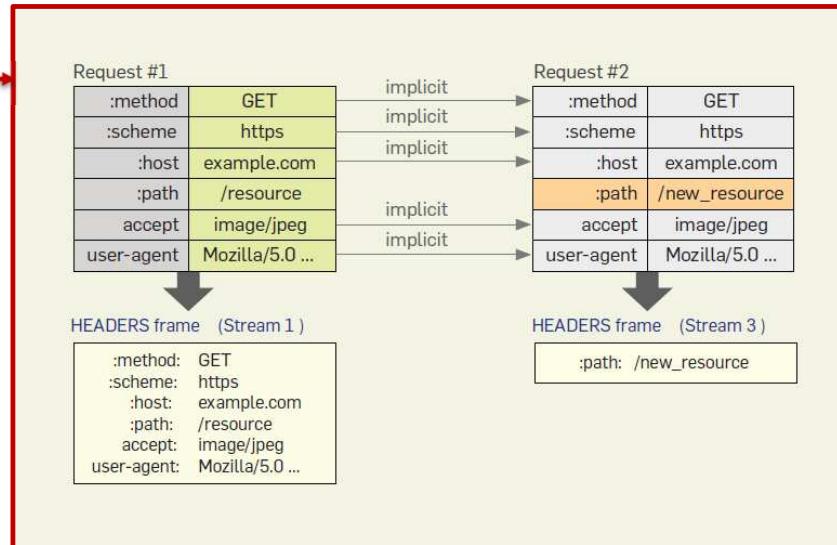
# HTTP/2.0

- Binary Framing
  - More compact representation
  - Easier and more efficient to process
- Separation of data and control planes

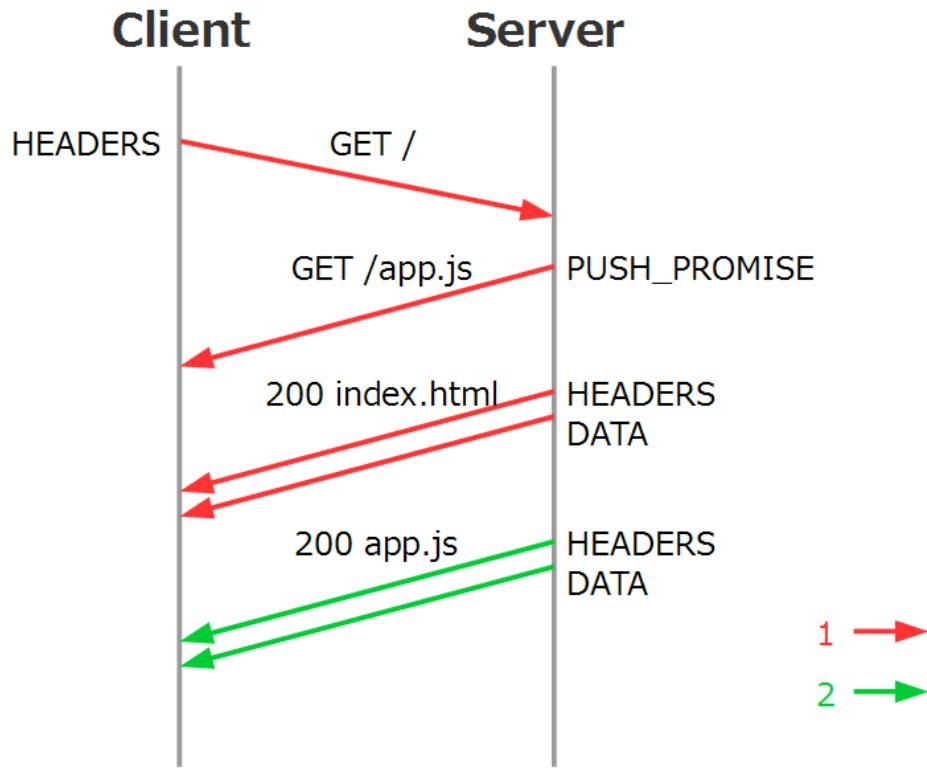


# HTTP/2.0

- Stream prioritisation
- Server push
  - Send several replies to a single request
- Header compression
- Flow control support
- HTTP 2.0 upgrade and discovery



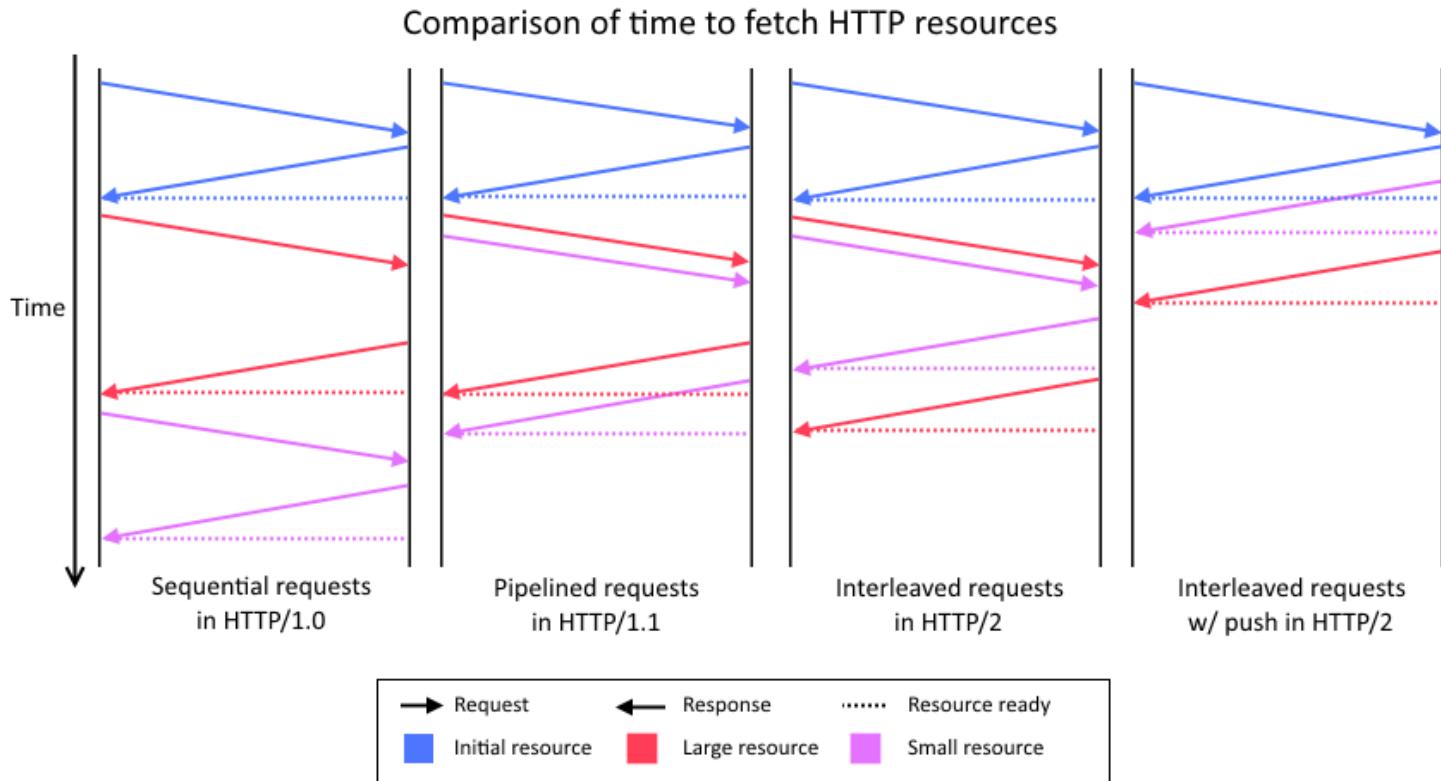
# GET w Multiple Elements



Source: <http://jinjor-labo.hatenablog.com/entry/2015/02/27/231753>

A. Aguiar

# HTTP/2.0

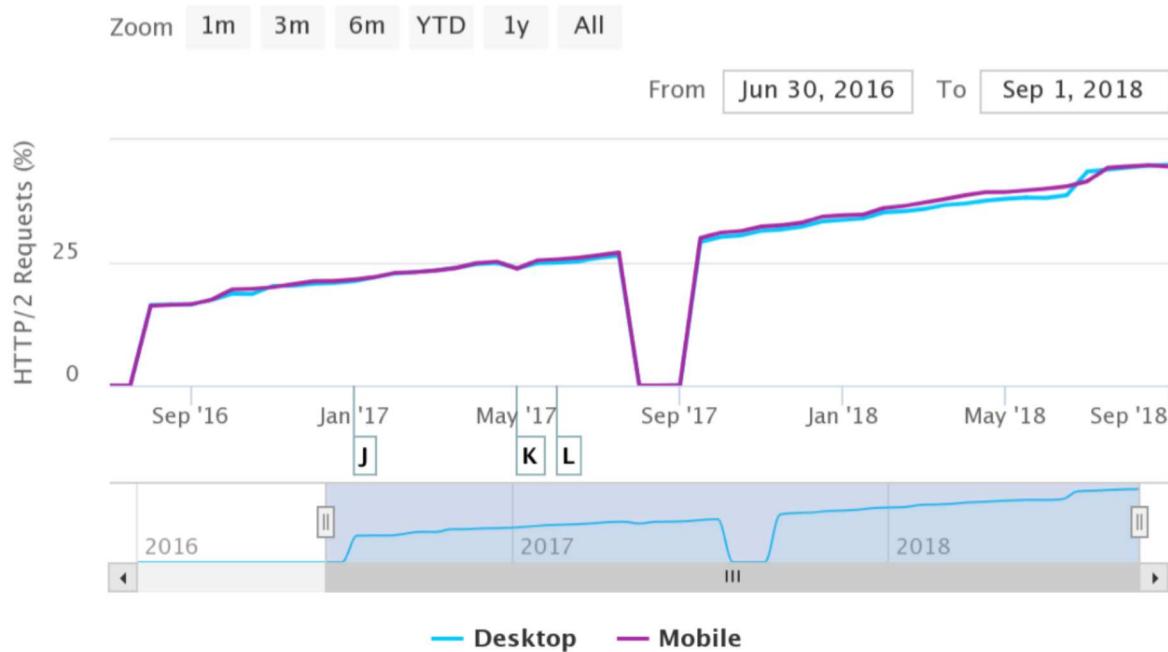


From [blog.scottlogic.com](http://blog.scottlogic.com)

# HTTP/2

Timeseries of HTTP/2 Requests

Source: httparchive.org



Next

## **CACHING, CONTENT DELIVERY NETWORKS**