

# SCOM/ SRSI

## Quality of Service

Ana Aguiar  
DEEC, FEUP  
2018-19

---

# Contents

- Resource Allocation Mechanisms
    - Queueing, dropping
    - Scheduling
    - Traffic shaping: leaky bucket, token bucket
  - QoS service models: IntServ and DiffServ
  - MPLS
-

# Resource Allocation

- Resources: links, routers
  - Resources shared among different packet flows
  - Resource allocation deals with mechanisms to share available resources
    - According to service models
    - Under normal operation and in case of high load
-

# Best Effort Networks

- Packets are treated individually
    - No notion of flow
  - Best-effort networks treat all packets alike
    - No service differentiation
  - Network and hosts do not “talk” to each other
  - Need to react to or avoid congestion
    - Congestion control and avoidance mechanisms
-

# Real-Time Communication

- Communication of audiovisual and multimedia data requires time correlation between sender and receiver
    - Guaranteeing throughput is not enough
  - Other applications
    - Control applications: robots, automation, IoT
    - Emergency, safety (e.g. autonomous driving)
    - Interactive data applications, e.g. trading
-

# Real-Time Communication

- Real-time applications require **delay guarantees** from the network
    - Best-effort networks offer no guarantees
  - Hence the need for other service models that can co-exist with best-effort
    - Packets are treated differently, according to the service they require
    - Required service may not always be supported
-

# What is Quality of Service (QoS)?

- Capability of the network to offer service guarantees different from best-effort
  - Capability of the network to provide different types of service
  - Enabling hosts to request service guarantees
-

# QoS Components

- QoS-capable network should provide
    - Abstraction of application requirements
    - Communication between hosts/ apps and network
    - Mechanisms for implementing service guarantees
      - Admission control
      - Policing of admitted flows
      - Resource allocation
      - Stateful routers
-




# Resource Allocation Implementation

- A few important mechanisms to implement QoS
    - Queueing and dropping discipline
      - How to manage queues at routers and treat packets when queues start to fill
    - Traffic Shaping
      - How to guarantee traffic behaviour
    - Scheduling
      - How to implement a specific service at a router
-

# Dropping Discipline


- If queue is full and a new packet arrives, which packet will be dropped?
  - Newest (tail),
  - Oldest (head),
  - Lowest priority
  - Random, ...
- Drop before filling up




# Queueing Discipline

- Which packet in the queue will be served next?
  - Oldest (FIFO)
  - Newest (FILO)
  - Highest priority
  - Depending on deadline
  - ...
- Multiple queues


First-in,  
First-out



First-in,  
Last-out




Independent of arrival




# Scheduling: Priority Queueing

- Packets can have different priorities
- Higher priority packets are served first
  - Low priority packets can be starved
- We have other options



# Scheduling: Fair Queueing

- Assume fair = equal
- One queue per flow
- Round-robin scheduling implements fair access (= transmission opportunity) to shared resource



# Scheduling: Fair Queueing

- Fair queueing is a scheduling discipline
- Not all packets have the same length
  - Round robin is not fair in this case
- Simulate bit-by-bit service using a **virtual clock**

$$F_i = \max(F_{i-1}, A_i) + P_i$$

$F_{i-1}$  finish time of packet  $i - 1$

$A_i$  arrival time of packet  $i$

$P_i$  time necessary to transmit packet  $i$

- Choose queue to serve according to simulated packet virtual finish time
-

# Scheduling: Fair Queueing (FQ) Universidade do Porto FEUP Faculdade de Engenharia

- If  $n$  flows have data to send, each flow receives  $1/n$  th of the bandwidth
    - Link is equally shared among backlogged flows
  - If only one flow has data to transmit, it receives the whole bandwidth
    - Work conserving scheduling policy: as long as there is data to transmit, the link is not idle
-

# Weighted Fair Queueing (WFQ) Universidade do Porto FEUP Faculdade de Engenharia

- Extends FQ to support weights per queue
  - Each queue receives service according to its weight
    - Instead of transmitting  $1/n$  bits of each queue per clock tick, it transmits  $w_i/\text{sum}(w_i)$
    - FQ is a special case of WFQ, where  $w_i=1$  for all  $i$
-



# Weighted Fair Queueing

- Enables the implementation of concepts of fairness other than fair  $\Leftrightarrow$  equal
- But does not specify how to do it
  - How to set weights is policy, not scheduling
  - Weight of a flow does not guarantee a minimum bandwidth to that flow, only a minimum percentage of available bandwidth

Policy  $\neq$  Mechanism that implements it

---

# Exercise

- 3 input flows, 1 output
- Initially, all queues empty
- Name the transmission order for
  - Fair queueing
  - Weighted fair queueing for  $w_1=2$ ,  $w_2=w_3=1$


Packet	Size	Flow
1	100	1
2	100	1
3	100	1
4	100	1
5	190	2
6	200	2
7	110	3
8	50	3

# Traffic Policing and Shaping

- Mechanisms:
    - Leaky bucket
    - Token bucket
  - Shaping: control traffic entering a network
  - Policing: identify traffic that does not conform
  - Marking: mark packets that do not conform to certain shape
-


# Leaky Bucket

- Single server queue
- Output data rate is constant
- Queue determines bursts that can be accommodated without losses
  - But output data rate is still constant
- Queue full => data loss




# Token Bucket

- Tokens are inserted into bucket at a rate
- Data transmission consumes tokens from bucket
  - Average output data rate equals token rate




# Token Bucket

- Data can only be transmitted for which enough tokens exist in the bucket



# Token Bucket

- Size of bucket indicates maximum burst size



# Token Bucket: Example

- Flow A
  - average rate = 1 Mbps
  - no variations
- Flow B
  - average rate = 1 Mbps
  - maximum bursts of 2 Mbps for at most 1 s, max 1/3 of the time




Image in "Computer Networks: a Systems Approach",  
L. Peterson, B. Davie



# Token Bucket: Example

- Flow A
  - $r_A = 1$  Mbps
  - $B_A = \text{MSS Bytes}$
- Flow B
  - $r_B = 1$  Mbps
  - $B_B = 1$  Mbit
  - Use only half the tokens during first 2 seconds
  - Use saved tokens during burst




Image in "Computer Networks: a Systems Approach",  
L. Peterson, B. Davie

# Exercise

- Assume processing rate  $R = 1$  packet (of length  $P$ ) per time unit.
  - The input flow has a burst of 10 packets, followed by an idle period of 20 packet times (totally there are 30 packet times).
  - Leaky bucket and Token bucket are described with the usual parameters where the bucket size  $B = 5 \times P$ .
    - What is maximum permissible burst without loss for the leaky bucket? What is the maximum delay for a packet?
    - Show the outputs of the flow shaped by a leaky bucket and a token bucket if the burst arrival of the 10 back packets observe the maximum permissible burst limit. Note the processing rate of a token bucket can be adjusted by the token generation rate.
    - If the maximum permissible burst limit is not followed by the leaky bucket algorithm, what will happen? How about the token bucket algorithm?
-

# QOS SERVICE MODELS

---

# QoS Paradigms

- Support requirements per flow  
Fine-grained QoS support  
Integrated Services (IntServ) and ATM
  - Support requirements per class of flows  
Coarse QoS support  
Differentiated Services (DiffServ)
-

# Integrated Services (IntServ)

- RFC 1633, 1994
  - Service guarantees per flow
  - Service based on end-to-end reservation
  - Specification of
    - service classes
    - reservation protocol (RSVP)
    - mechanisms to implement reservation
-

# Service Classes

- Guaranteed service (RFC 2211)
    - Guaranteed maximum delay per packet
    - Designed for intolerant applications
  - Controlled load (RFC 2212)
    - Designed for applications that tolerate some loss
    - Provides flow isolation so that flows perceive network as lightly loaded
  - Best Effort
-

# Resource ReSerVation Protocol Universidade do Porto FEUP Faculdade de Engenharia

- RFC 2205
  - Signalling protocol for resource reservation
  - Provides
    - means for hosts to describe traffic and requirements to intermediate nodes
    - means for routers to inform hosts whether they have enough available resources
-

- Specifies flow traffic (TSpec) and flow requirements (RSpec)
  - TSpec
    - Used in admission control calculations by the network
    - Should describe traffic as precisely as possible
    - Most flows have varying traffic patterns
      - E.g., recall variations in data rate of coded video
    - Network must accommodate simultaneous bursts
  - RSpec
    - Guaranteed service: target or maximum delay
    - Controlled load: requires no additional specification
  - Sent using Resource Reservation Protocol (RSVP)
-



# Traffic Description

- Traffic description should be very precise
  - Token bucket describes traffic
    - If a flow does not expect bursts, it should advertise a very small bucket size
    - Resources are reserved according to TSpec and RSpec
-

# RSVP: Resource Reservation

- Sender sends TSpec to receiver in PATH message
  - Each router saves the reverse path for the PATH message to use later in reverse-forwarding the reservation message
-

# RSVP: Resource Reservation

- Receiver sends RESV message along reverse path, containing TSpec and its own RSpec
  - Each router assesses whether it can provide the reservation (admission control)
    - RESV message is sent to previous router in path if yes
    - Error message is sent to requesting receiver if not
-

# RSVP Resource Reservation

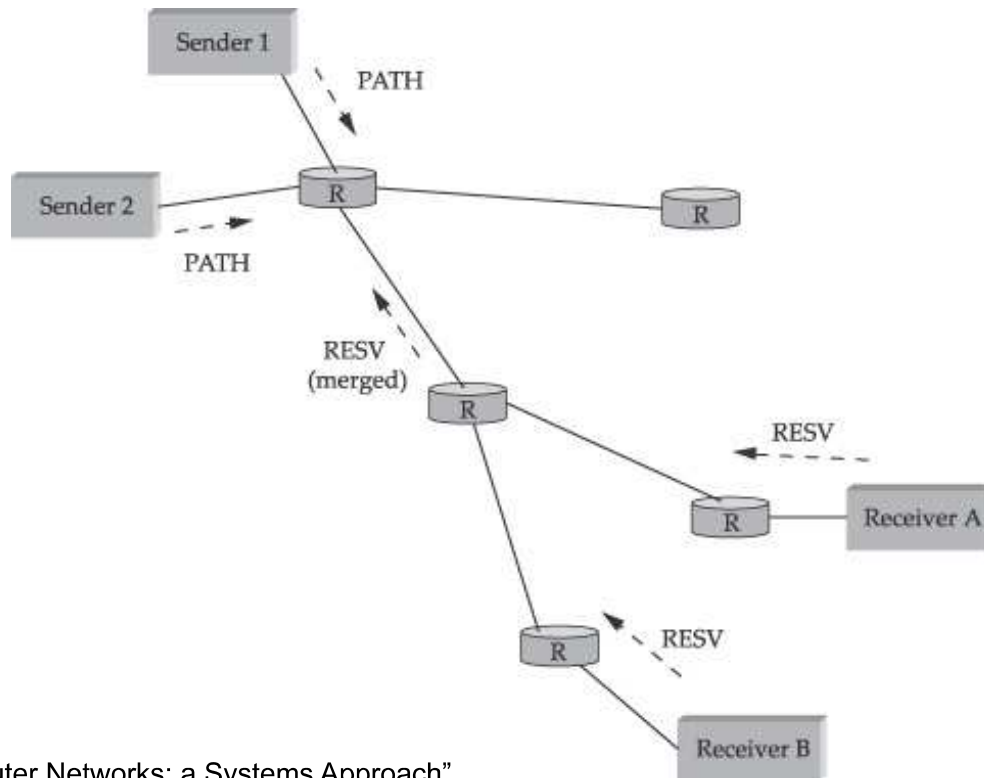


Image in "Computer Networks: a Systems Approach",  
L. Peterson, B. Davie

# RSVP: Soft state

- Routers keep per flow soft state
    - state must be refreshed periodically to be kept alive
    - state is discarded in absence of refresh message
    - maintains robustness of connectionless network in case of hosts or routers crashing or rebooting
  - Need to periodically refresh state makes adaptation of reserved resource easy
-

# RSVP: Route Changes

- PATH message is re-sent every 30 s
    - Recall: RESV messages also periodically re-sent
  - Routers adapt their reverse path according to new PATH message
  - If route has changed, next RESV message will be passed along new route and trigger reservation along it
-

# RSVP: Multicast Support

- Multicast
    - More receivers than senders expected
    - Receivers may have different requirements
  - Reservation requests are started by receiver
-

# RSVP: Multicast Reservation

- Multicast trees have common path segments to different receivers
  - RESV messages are only forwarded if existing reservations are not sufficient
  - Otherwise, RESV messages from different receivers are merged
-



# Admission Control

- Mechanism to decide whether a new flow can be accepted
  - Both Guaranteed Service and Controlled Load require admission control
  - Admission of a new flow should not cause any on-going flow to receive less than requested service
    - Decision depends on FlowSpec and available resources
-

- Guaranteeing that an accepted flow does not exceed his TSpec
    - Misbehaviour can lead to not meeting service reservations for other flows
    - Packets out of specification can be always dropped or forwarded anyway as long as they do not interfere with other flows
    - Can you think of a mechanism to implement this?
-

# Packet Classification

- Delivering different service to different packets requires packet classification
  - Use source and destination addresses, ports, and protocol number field
    - These fields must be parsed at each router
-

# Policy

- Policies are sets of rules that define the behaviour of mechanisms
    - For example, flows from certain hosts (e.g. plant controller, video conference system) have higher priority, or can preempt others in admission control
-

# Packet Scheduling

- Scheduling is not specified in the architecture or service model
  - IntServ is commonly associated with WFQ
    - Each guaranteed service flow has own queue
    - Controlled load services shared one queue
  - Implementing different services in a node is very complex and scales badly
-

# IntServ: Scalability

- IntServ has drawbacks
  - Routers must
    - calculate admission control for flows
    - classify packets
    - manage different queues and a scheduling algorithm
    - maintain state for all flows
  - Reservation maintenance messages are periodically sent by each sender and receiver
-

# IntServ: Scalability


- Take a 10 Gbps link and suppose it is filled up with 64 kbps audio streams
  - 156 250 flows
    - To keep state, classify packets, admit or decline
    - And all the other functionalities mentioned
  - IntServ is often considered unscalable
    - Is it nowadays still so?
-

# Differentiated Services

- Provides coarser quality of service inside one administrative domain (AD)
  - Few pre-defined QoS classes
  - Flows aggregated according to service class
  - Each packet carries the indication of which service class it belongs to
  - Service provided per hop
  - Good fit to operator business model
-



# DiffServ in a Domain



# DiffServ Architecture

- Service agreement regarding traffic aggregates
    - Service Level Agreement (SLA) between ISP and clients
  - Edge routers do traffic conditioning
    - Classify and mark packets according to SLA
    - Do traffic policing (ingress) and shaping (egress)
  - Domain routers implement per-hop behaviour
    - Queueing and scheduling
-

# DiffServ: Per Hop Behaviour

- DiffServ implements service differentiation through different per-hop behaviours
  - PHB defined for each packet: fully stateless
  - Use header field to differentiate packets
    - Re-define Type of Service (TOS) byte in IP header
    - 6-bit DiffServ Code Point (DSCP) field defines the per hop behaviour
-

# Expedited Forwarding (EF) PHB Universidade do Porto FEUP Faculdade de Engenharia

- Packets forwarded as soon as possible
    - Building block for low loss, low delay, low jitter
  - Rate of EF packets arriving at a router should be lower than router capacity
  - Requires some form of admission control
    - At edge router, before marking packets as EF
    - Example admission control  $\sum \lambda_{EF} < \min_i \{r_i\}$ ,  $r_i$  is the rate of link  $i$  in the network
-

# Expedited Forwarding (EF) PHB Universidade do Porto FEUP Faculdade de Engenharia

- Implemented by queueing policies
    - Priority queueing with strict priority for EF and AF served with WFQ
  - OR
  - Weighted Fair Queueing with much larger weight for EF queues
    - WFQ can guarantee that other traffic will not be starved at the cost of slight service inaccuracy
    - EF traffic may not always get lowest possible delay
-

# Assured Forwarding (AF) PHB

- AF PHB offers services that are extensions to RED
    - RED In and Out (RIO)
    - Weighted RED (WRED)
-

# DiffServ Building Blocks

Ingress routers

- Classification
- Policing

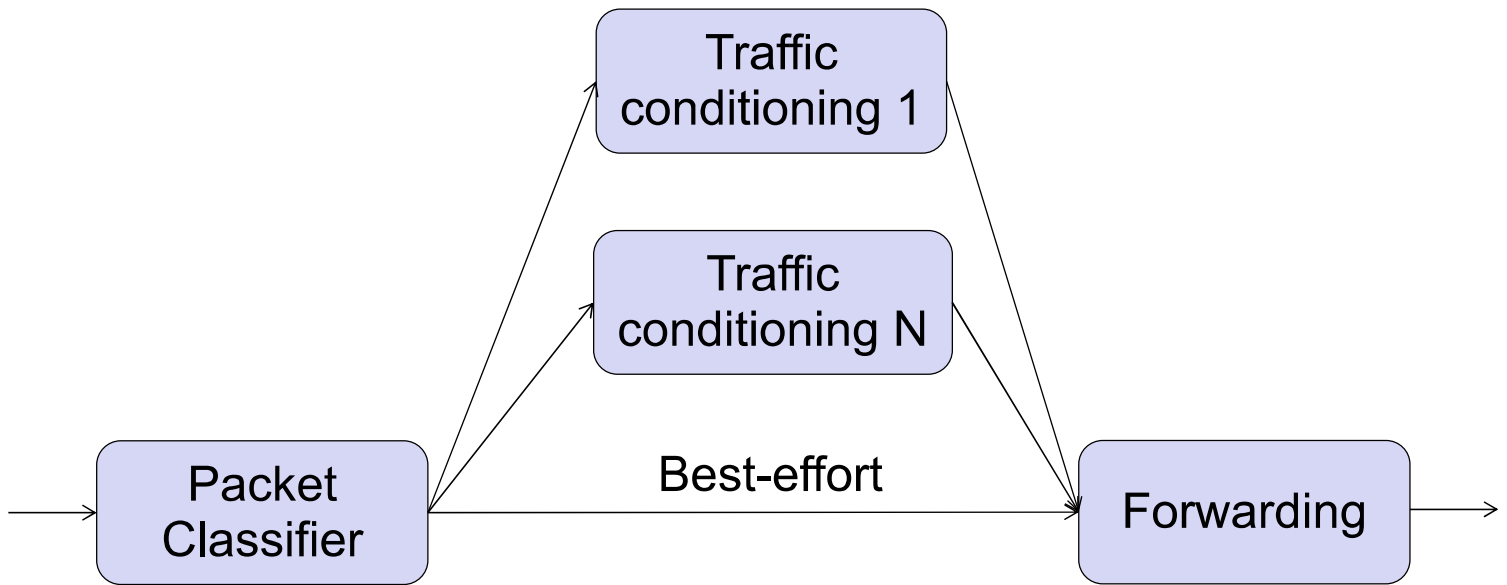
Domain routers

- Flow isolation
- Queueing
- Scheduling

Egress routers


- Traffic shaping
-

# Ingress Edge Router





# Traffic Conditioning




# Classification

- Classifier at ingress router identifies and marks packets according to the service class
  - Packets can be classified based on
    - sending domain
    - transport protocol
    - ... (header fields)
-

- At ingress router, measure amount of traffic arriving for each flow aggregate
    - Use token buckets
  - Mark packets as in or out of profile
    - Routers in the network can decide whether they can handle out of profile packets or not
    - Enables better resource sharing than dropping
    - RIO, WRED (later on)
-

# Domain Router



# Flow Isolation


- Service delivered to one class should be independent of service delivered to another
  - Packets in different service classes are placed in different queues at each router
-

# DiffServ with Priority Queueing

- EF packets have highest priority
  - AF has several queues with different service
  - If packets of an aggregate use different queues, packet re-ordering can happen
-

# RED In and Out (RIO)

- In or Out refers to the traffic profile agreed with the sender
  - Incoming traffic will be marked as In or Out of profile
  - Traffic exceeding agreed volume will follow Out curve drop probability, and be dropped earlier at the domain routers



# RIO Considerations

- RIO will try to keep congestion low so that In profile packets are rarely dropped
  - Nevertheless, performance depends on overall bandwidth utilisation
    - AF and EF co-exist in the domain
  - Performance depends on parameterisation and admission control
-



# Weighted RED

- More than two drop probability curves can provide more than two classes of service



# DiffServ with WFQ

- Different weights are used to provide different levels of service
  - Difficulty lies in setting the queue weights
    - Queue weights for premium traffic can be set very high, but there is a certain margin of error in service
  - WFQ can be combined with WRED
    - Packets in each queue can have different drop probability curves
-

# DiffServ Example

- EF class whose packets are forwarded ASAP
    - Because of this, only in-profile EF packets can be admitted into the domain
  - 4 AF classes with 3 drop levels each
  - 1 Best Effort class
  
  - No standard for bandwidth allocation
  - No packet re-ordering within a class
-

# Traffic Shaping

- At egress routers
    - Traffic arriving from network is bursty and irregular
    - Leaves domain shaped according to leaky bucket
      - Regular data rate
      - Queue size determines acceptable burstsize/ dropping
  - Guarantees that domain is “well-behaved” towards others
-

# EF vs AF PHB

How would you classify the discussed applications?

VoIP, VoD, WWW, file transfer, control...

---

# Admission Control

- Before accepting a flow for a requested service, verify that the network can deliver that service
    - IntServ and EF PHB require admission control
  - How aggressive should it be?
    - Too conservative can lead to low resource usage
    - Too aggressive can lead to unfulfillment of promised service agreements
    - Depends on the goal of running a network
      - Serve as many flows as possible?
      - In terms of throughput or delay?
    - Maximise total throughput? Minimise delay?
    - ...
-

# DiffServ vs IntServ

- Type of service that can be delivered
  - Granularity
  - Setup of the service
  - Service based on reservation vs provisioning
  - Scope of service: local vs end-to-end
-

# MULTI-PROTOCOL LABEL SWITCHING

---



# Multi-protocol Label Switching

- Multi-protocol Label Switching (MPLS)
    - Between link and network layer
    - Supports multiple protocols underneath and above
  - Labels (“shim”) are added to packets according to Forwarding Equivalence Class (FEC)
  - Label Switching Routers (LSR) forward according to label
-

# MPLS

- Labels can be added to packets or flows
  - Packets can have more than 1 label
    - To enable multiple routing levels
  - LSRs forward according to “last in” label
  - LSRs can swap labels
  - Class of Service can be inferred from the label
-

- Hop-by-hop routing
    - Regular routing, e.g. OSPF
    - LSRs create labels and distribute them among peers
    - LDP: Label Distribution Protocol
  - Explicit routing
    - Ensures that a certain FEC follows a pre-determined path
    - Use reservation protocol to distribute labels (RSVP-TE)
      - Label allocated by receiving node
      -
-

# **NEXT WEEK: NETWORK NEUTRALITY**

---